

FactorJoin: A New Cardinality Estimation Framework for Join Queries

ZINIU WU, Massachusetts Institute of Technology, USA

PARIMARJAN NEGI, Massachusetts Institute of Technology, USA

MOHAMMAD ALIZADEH, Massachusetts Institute of Technology, USA

TIM KRASKA, Massachusetts Institute of Technology, USA

SAMUEL MADDEN, Massachusetts Institute of Technology, USA

Cardinality estimation is one of the most fundamental and challenging problems in query optimization. Neither classical nor learning-based methods yield satisfactory performance when estimating the cardinality of the join queries. They either rely on simplified assumptions leading to ineffective cardinality estimates or build large models to understand the complicated data distributions, leading to long planning times and a lack of generalizability across queries.

In this paper, we propose a new framework *FactorJoin* for estimating join queries. *FactorJoin* combines the idea behind the classical *join-histogram* method to efficiently handle joins with the learning-based methods to accurately capture attribute correlation. Specifically, *FactorJoin* scans every table in a DB and builds single-table conditional distributions during an offline preparation phase. When a join query comes, *FactorJoin* translates it into a factor graph model over the learned distributions to effectively and efficiently estimate its cardinality.

Unlike existing learning-based methods, *FactorJoin* does not need to de-normalize joins upfront or require executed query workloads to train the model. Since it only relies on single-table statistics, *FactorJoin* has a small space overhead and is extremely easy to train and maintain. In our evaluation, *FactorJoin* can produce more effective estimates than the previous state-of-the-art learning-based methods, with 40x less estimation latency, 100x smaller model size, and 100x faster training speed at comparable or better accuracy. In addition, *FactorJoin* can estimate 10,000 sub-plan queries within one second to optimize the query plan, which is very close to the traditional cardinality estimators in commercial DBMS.

CCS Concepts: • **Information systems** → **Query optimization**.

Additional Key Words and Phrases: Query Optimization, Cardinality Estimation

ACM Reference Format:

Ziniu Wu, Parimarjan Negi, Mohammad Alizadeh, Tim Kraska, and Samuel Madden. 2023. FactorJoin: A New Cardinality Estimation Framework for Join Queries. *Proc. ACM Manag. Data* 1, 1, Article 41 (May 2023), 27 pages. <https://doi.org/10.1145/3588721>

1 INTRODUCTION

Cardinality estimation (CardEst) is a critical component of modern database query optimizers. The goal of CardEst is to estimate the result size of each query operator (i.e., filters and joins), allowing

Authors' addresses: Ziniu Wu, ziniu@mit.edu, Massachusetts Institute of Technology, Cambridge, MA, USA, 02139; Parimarjan Negi, pnegi@mit.edu, Massachusetts Institute of Technology, Cambridge, MA, USA, 02139; Mohammad Alizadeh, alizadeh@mit.edu, Massachusetts Institute of Technology, Cambridge, MA, USA, 02139; Tim Kraska, kraska@mit.edu, Massachusetts Institute of Technology, Cambridge, MA, USA, 02139; Samuel Madden, madden@mit.edu, Massachusetts Institute of Technology, Cambridge, MA, USA, 02139.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2023 Copyright held by the owner/author(s). Publication rights licensed to ACM.

2836-6573/2023/5-ART41

<https://doi.org/10.1145/3588721>




Fig. 1. Existing CardEst approaches for handling join queries. (a) The Selinger model writes the cardinality of a two-table join query as the product of filter selectivity P_A, P_B on both tables and the estimated join size $|A \bowtie B|$. Using join-key uniformity, this is estimated as the product of table sizes $|A| * |B|$ divided by the maximal number of unique values of the join keys. (b) The join-histogram bin the domain of join keys Id_1, Id_2 as histograms, assumes join-key uniformity within each bin. (c) The learned data-driven methods de-normalize $J = A \bowtie B$ and learn the distribution P_J to estimate the cardinality of Q . (d) Our framework captures the correlation between join keys and filter predicates to accurately estimate $|Q|$ without de-normalization.

the optimizer to select the most efficient join ordering and physical operator implementations. An ideal CardEst method satisfies several properties: it is *effective* at generating high-quality query plans, *efficient* so that it minimizes estimation latency, and easy to *deploy* in that it has a small model size, fast training times, and the ability to scale with the number of tables and generalize to new queries. Unfortunately, existing CardEst techniques do not satisfy at least one of these properties. The reason behind these failures can largely be reduced to the handling of two related problems: (1) how to characterize attribute correlations and (2) how to model the distribution of join-keys, which directly determine the size of joins.

Background: Many deployed CardEst approaches are still based on the Selinger model [60] and assume *attribute independence*, where the model ignores correlations between attributes, and *join-key uniformity*, where the model further assumes that join keys have uniformly distributed values. These assumptions provide a simple way of estimating the cardinality of join queries using only single-table statistics (e.g., histograms) over single attributes, which are easy to create and maintain. As a result, the Selinger-based CardEst techniques are extremely efficient and easy to deploy but at the cost of effectiveness because most real-world datasets contain complex attribute correlations and skewed join-key distributions.

To overcome the limitations of these assumptions, many commercial database systems try to relax them to increase their effectiveness. For example, past works [7, 26, 29] propose and some systems (e.g. Oracle [8]) employ *join-histograms* that relax the *join-key uniformity assumption*. Specifically, these methods build frequency histograms on the join keys and then construct the join distribution by “multiplying” the histograms to estimate the join size (see Figure 1(b)). In this way, these methods can more precisely capture the join-key distributions to generate more effective estimates of the join cardinality. However, they still assume that the join keys are uniformly distributed within each bin of the histogram.

Similarly, there have been works on relaxing the *attribute independence assumption*. For example, multi-dimensional histograms [9, 14, 15, 43, 52, 66], self-tuning histograms [4, 12, 31, 61], and

singular value decomposition [58], have been proposed to capture the correlation between attributes. However, their effectiveness and/or efficiency are still not satisfactory [18, 75]. More recently, learning-based methods which learn the underlying data distributions [23, 33, 75, 78] have been proposed to more accurately and compactly represent the attribute correlations within a single table, but these still have efficiency concerns that we detail below.

Specifically, the learned data-driven methods [13, 25, 64, 71, 74, 78] analyze data and build distributions for all join patterns in a database. They need to de-normalize the joined tables and add a potentially exponential number of extra columns. Then, they build distributions over the de-normalized tables to characterize all attribute correlations and handle joins (see also Figure 1(c)). This allows the learned data-driven methods to be highly accurate for join estimates at the cost of slow training time and large model size (i.e., worse deployability). Alternatively, learned query-driven methods [33, 42, 63] circumvent the *join-key uniformity* assumption by building supervised models to map the join queries to their cardinalities. However, they require an impractical number of executed queries to train their models, which is unavailable to new DB instances and do not generalize well to unseen queries.

Our Approach: Surprisingly, we are not aware of any work which tries to combine the advances of learning-based methods for accurately characterizing correlations with the idea of *join-histograms*. In this work, we develop a novel CardEst framework called *FactorJoin* that combines these two approaches, i.e., using *join-histograms* to efficiently handle joins coupled with learning-based methods to accurately capture attribute correlation. This is not a straightforward extension. As we will show, adapting the idea of *join-histograms* while preserving the benefits of using learning-based methods for attribute correlations is non-trivial. *FactorJoin* only builds sophisticated models to capture the correlations within a single table. The choice of the model to capture correlation is orthogonal to the techniques of *FactorJoin*, however, we do require the model to be able to provide conditional distributions of one or more keys. In our current implementation of *FactorJoin*, we use sampling [41] and Bayesian Networks [34] as they are fast to train and execute.

These single-table conditional distributions are then combined using a *factor graph model* [44]. Specifically, *FactorJoin* translates a join query Q into a factor graph over single-table data distributions. This allows us to formulate the problem of estimating the cardinality of Q as a well-studied inference problem [34, 37, 44] on this factor graph. Conceptually, there exist some similarities to the *join-histogram* technique. However, our formulation allows *FactorJoin* to generalize to cyclic and self-joins. It also improves the computation efficiency by relying on existing inference algorithms [35, 46], and facilitates scaling to arbitrary join sizes. Finally, in contrast to the idea of *join-histogram* and many other techniques in the space, we propose to actually not estimate the expected join cardinality, but rather use a probabilistic upper bound. As shown by others [1, 3, 5, 22], under-estimation is often worse than over-estimation, and tailoring the estimate to a probabilistic upper bound can significantly improve the end-to-end query performance.

As a result, *FactorJoin* is able to estimate the cardinality of join queries using only single table statistics with similar or better estimation effectiveness than the state-of-the-art (SOTA) learning-based techniques. When compared with the learned data-driven methods, *FactorJoin* has a much smaller model size, faster training/updating speed, and easier for system deployment because we do not need to denormalize the join tables or add exponentially many extra columns. In addition, *FactorJoin* supports all forms of equi-joins, including cyclic joins and self joins, as well as complex base table filter predicates, including disjunctive filter clauses and string pattern matching predicates because it can flexibly plug-in various types of base-table estimators to support them. These queries are not supported by the existing learned data-driven methods [25, 74, 78], which require the join template to be a tree. Unlikely the learned query-driven methods [33, 42, 63], *FactorJoin* does not

depend on the executed query workload, so it is robust against data updates and workload shifts, can be quickly adapted to a new DB instance, and is generalizable to new queries. However, in presence of query workload, *FactorJoin* can incorporate this information to further optimize the model construction process.

We integrate *FactorJoin* into Postgres' query optimizer and evaluate its end-to-end query performance on two well-established and challenging real-world CardEst benchmarks: STATS-CEB [18] and IMDB-JOB [38]. On the STATS-CEB, *FactorJoin* achieves near-optimal performance in end-to-end query time. It has comparable performance to the previous SOTA learned data-driven model (FLAT [78]) in terms of estimation effectiveness but with 40x lower estimation latency, 100x smaller model sizes, and 100x faster training times. On the IMDB-JOB, our framework achieves the SOTA performance in end-to-end query time (query execution plus planning time). Specifically, the pure execution time of our framework is comparable to the previous SOTA method (pessimistic estimator [5]), but our estimation latency is 100x lower, making our overall end-to-end query time significantly faster. Specifically, our framework can estimate 10,000 sub-plan queries in one second to optimize the query plan, which is close to the planning time of the traditional CardEst method used in Postgres. We also carry out a series of controlled ablation studies to demonstrate the robustness and advantages of different technical novelties of *FactorJoin*.

In summary, our main contributions are:

- We formulate the problem of estimating the cardinality of join queries in its general form as a factor graph inference problem involving only single-table data distributions (Section 3).
- We propose a new binning and upper bound-based algorithm to approximate the factor graph inference (Section 4).
- We design attribute causal relation exploration and progressive estimation of sub-plan queries techniques to improve the efficiency of our framework (Section 5).
- We conduct extensive experiments to show the advantages of our framework (Section 6).

Before describing the details of *FactorJoin*, we begin with a detailed description of existing CardEst approaches.

2 BACKGROUND AND ANALYSIS

In this section, we first define the CardEst problem and then analyze existing CardEst methods and how our approach differs from them.

2.1 CardEst problem definition

Single table CardEst: Let A be a table with n attributes $A_1, \dots, A_n$. A single table selection query $Q(A)$ on A can be viewed as a conjunction or disjunction of filter predicates over each attribute, e.g. $Q(A) = (A_1 \in [0, 5]) \vee (A_2 \text{ LIKE } \%An\%) \wedge (A_5 \leq -10)$.

Let $|Q(A)|$ denote the cardinality of $Q(A)$, i.e. the number of tuples in A satisfying $Q(A)$. Let $P_A(Q(A))$ denotes the selectivity (probability) of tuples satisfying $Q(A)$. Then, we have $|Q(A)| = P_A(Q(A)) * |A|$, where $|A|$ denotes total number of tuples in A .

CardEst for multi-table join queries: Consider a database instance $\mathcal{D}$ with m tables $A, B, \dots, M$, a query Q consists of a join graph representing the join conditions among the selected tables and a set of base-table filter predicates. For simplicity, we abuse the notation and use $Q(I)$ to denote the filter predicates of Q on table I .

Let Ω denote the denormalized table resulting from the join conditions of Q . Then, the cardinality $|Q| = P_\Omega(Q(A), Q(B), \dots, Q(M)) * |\Omega|$, where $Q(I)$ can be empty set if table I is not touched by Q . Estimating Q is particularly challenging because there can exponential number of join patterns in the database, each associated with a unique Ω and distinct probability distribution P_Ω . Therefore,

Cardinality Estimation Methods Characteristics		Traditional			Learned		Bound based [1, 3, 5, 22]	FactorJoin (Ours)
		Sampling [39, 40, 41, 76]	Selinger model [10, 59, 60]	Join Histogram [7, 26, 29]	Data-driven [13, 23, 64, 65, 67, 71, 74, 78]	Query-driven [33, 42, 54, 63]		
Performance	Effective	—	—	—	✓	✗	✓	✓
	Efficient	—	✓	✓	—	✓	—	✓
	Deployment	Small model size	✓	✓	—	✓	✓	✓
		Fast train/update	✓	✓	—	—	✓	✓
		Scalable with no. of joins	—	✓	—	✓	✓	✓
		Generalizes to new queries	✓	✓	✓	—	✓	✓
		Supports cyclic join	✓	✓	—	✓	✓	✓
Techniques	Handling correlation	Use Sampling	✓	—	—	—	—	✓
		Use ML models	—	—	—	✓	—	✓
	Handling Join	Uses binning/histograms	—	—	✓	✓	—	✓
		Denormalizes join tables	—	—	✓	—	—	✓
		Adds extra columns	—	—	✓	—	—	—
		Uses query information	—	—	—	✓	✓	—
		Uses bound	—	—	—	—	✓	✓

Table 1. Summary of existing CardEst methods.

to accurately estimate different join queries, the CardEst methods need to capture exponential many complicated data distributions.

2.2 Analysis of existing CardEst methods

Since the focus of this paper is multi-table join queries, we defer the discussion of single-table estimators to Section 7 and focus on analyzing different approaches for handling join queries. We can categorize these approaches into four classes: *traditional*, *learned data-driven*, *learned query-driven*, and *bound-based*. We summarized their characteristics in Table 1 and provide their details as follows.

Traditional methods: The sampling-based methods [39–41, 77] construct a small sample on each table and join these samples to estimate the cardinality of join queries. The traditional histogram-based methods generally adopt the attribute independence and join uniformity assumptions to decompose the join queries as a combination of single table estimates. Taking a query Q joining two tables A and B as an example, these assumptions will simplify the cardinality $|Q| = P_{\Omega}(Q(A), Q(B)) * |\Omega|$ as $P_A(Q(A)) * P_B(Q(B)) * |\Omega|$, where $P_A(Q(A))$ is the estimated single table selectivity. There exist two widely-used approaches to estimate the size for the denormalized join table $|\Omega|$. First,

the Selinger models [10, 59, 60] assume that join keys have uniformly distributed values, collect the number of distinct values (NDV) on the join keys from both tables, and estimate $|\Omega|$ as $|A| * |B| / \max\{NDV(A), NDV(B)\}$. Second, the *join-histogram* methods [7, 26, 29] first bin the domain of the join keys as histograms and assume join uniformity assumption for the values in each histogram bin. Then they apply the *distinct values* methods for each bin and sum over all bins.

Our work follows the convention of the *join-histogram* methods but avoids the simplifying assumptions, which can not be trivially achieved. Naive approaches such as building the full multi-dimensional histograms can resolve the attribute independence assumption but would introduce unaffordable storage overhead. Many follow-up works [27, 28, 30] explore different histogram strategies to reduce the error caused by the join uniformity assumption. They either take an impractical amount of time to construct the histograms or cannot provide high-quality estimates.

In short, the traditional methods combine the single-table estimates to estimate the join queries. This approach is efficient, easy to train, update and maintain, thus perfect for system deployment. However, their simplifying assumptions can generate erroneous estimates and poor query plans.

Learned data-driven methods: These methods circumvent the simplifying assumptions used in traditional methods by understanding the data distributions for the exponential number of join templates in a DB instance. Specifically, some methods [13, 64, 65] denormalize the join of all tables in a database and use Bayesian networks to model the distribution. The current SOTA approaches for handling join queries try to model the data distributions of the denormalized join tables by using fanout-based techniques [23, 67, 71, 74, 78]. These methods denormalize some tables and add a possibly exponential number of fanout columns, which are used to formulate the distributions for each join template. Using this method, the learned data-driven methods can produce effective estimations for join queries and high-quality query plans [18]. However, these models generally have a long training time, large model size, slow estimation speed, and unscalable performance with the number of joins in the query. Furthermore, the learned data-driven approaches currently can not handle self-joins or cyclic-joins and are very ineffective in processing complicated filter predicates such as string pattern matching queries.

Learned query-driven methods: These methods analyze the executed query workload, map the featurized join query directly to its cardinality using supervised models such as neural networks [33, 54], xgboost [11], tree-LSTM [63], and deep ensembles [42]. The query-driven methods generally have low estimation latency but the estimation effectiveness is highly dependent on the training workload; we put a semi-check mark (✓) for these methods in Table 1. To achieve reasonable estimation quality, they generally need an excessive amount of executed training queries [18], which are unavailable for new DB instances. Furthermore, the query-driven models need to retrain in the case of data update or query workload shift. An impractical amount of new executed queries are again needed for this re-training process.

Bound-based methods: Traditional CardEst methods in modern DBMS tend to severely underestimate the cardinality, thus sometimes choosing significantly more expensive query plans [3, 5, 55]. To avoid under-estimation, the bound-based methods [1, 3, 5, 22] use information theory to provide an upper bound on the cardinality. These bounds are very effective to help generate high-quality query plans as they can avoid expensive join orders and physical operators [18, 55]. However, these methods do not build single-table estimators to understand the data distributions. Instead, in presence of base-table filters, they need to materialize the filtered tables and populate the bound during query run-time, which can produce very high latency and overhead. Therefore, despite the bound-based methods providing meaningful insights into understanding joins, they can not be practically deployed in a DBMS.




Fig. 2. Example query Q joining two tables with condition $A.id = B.Aid$ and filter predicates $Q(A)$ and $Q(B)$.

Summary: None of the existing approaches can simultaneously satisfy the desired properties of CardEst, namely effective, efficient and appropriate for system deployment. However, each of them contains some advantageous techniques, inspiring us to build *FactorJoin* that can bridge the gaps amongst all these CardEst categories and unify their merits.

3 NEW FRAMEWORK FOR JOIN QUERIES

Given this background on previous CardEst approaches, we now introduce the high-level operation of *FactorJoin*. In Section 3.1, we elaborate on *FactorJoin*'s core technique: accurately estimating join queries with only single-table statistics using factor graph inference techniques. Then, in Section 3.2, we analyze the complexity of this inference procedure and motivate several optimizations in *FactorJoin*. Finally, we provide a workflow overview in Section 3.3.

3.1 Problem formulation

We first illustrate the main idea with a simple example of a two-table join query. Then, we formalize the join query estimation problem as a PGM inference problem over single table distributions.

Two-table join query example: Figure 2 illustrates a query Q joining tables A and B on the inner join¹ condition $A.id = B.Aid$ with base table filter predicates $Q(A)$ and $Q(B)$. Table A first goes through the filter $Q(A)$, resulting in an intermediate table $A|Q(A)$ (records in A that satisfy the filter $Q(A)$). The same procedure is applied to table B . Then, the query Q will match the value of join keys $A.id$ and $B.Aid$ from these two intermediate tables. Specifically, the value a appears 8 times in table $A|Q(A)$ and 6 times in table $B|Q(B)$, resulting in value a appearing 48 times in the join result. Therefore, we can calculate the cardinality of this query Q as $8 \times 6 + 4 \times 5 + 3 \times 5 = 83$.

We can formulate the above procedure for calculating Q as a statistical equation in Equation 1, where $D(A.id)$ denotes the domain of all unique values of $A.id$. We observe that only single-table distributions P_A and P_B are required to accurately compute the cardinality of this join query. Also in this equation, $P_A(A.id=v|Q(A)) * |Q(A)|$ equals to $P_A(A.id=v \wedge Q(A)) * |A|$, which is exactly what single-table CardEst methods estimate. Thus, we can *accurately* calculate the cardinalities of two-table join queries using single-table estimators. Note that the summation over the domain of join key $\mathcal{D}(A.id)$ has the same complexity as computing the join. Therefore, we need to approximate this calculation for *FactorJoin* to be practical. The details of our approximation schemes are given in Section 3.2 and Section 4.

$$|Q| = \sum_{v \in \mathcal{D}(A.id)} P_A(A.id = v|Q(A)) * |Q(A)| * P_B(B.Aid = v|Q(B)) * |Q(B)| \quad (1)$$

¹In this paper, we only discuss inner and equality joins. *FactorJoin* naturally supports left, right, and outer joins. We leave support for non-equal joins as future work.




Fig. 3. Formulating a general join query cardinality estimation problem into a factor graph inference problem. Figure (i) shows a SQL query Q_2 . Figure (ii) visualizes the join template of Q_2 . Figure (iii) presents the factor graph model to compute Q_2 's cardinality.

Formal problem formulation: A general join query can involve a combination of different forms of joins (e.g., chain, star, self, or cyclic joins) so its counterpart equation (as Equation 1) can be very difficult to derive and compute; we have put these equations/derivations in Section A.1 of our technical report [70]. We provide a generalizable formulation that automatically decomposes join queries into single-table estimations using the factor graph model.

Figure 3-(i) shows a SQL query Q_2 joining four tables A, B, C, D . We visualize its join template as a graph in Figure 3-(ii), where each dashed rectangle represents a table, each ellipse (node) represents a join-key in Q_2 , and each solid line (edge) represents an equi-join relation between two join keys connected by it. Note that both sides of an equi-join relation represent the semantically equivalent join keys. We call them *equivalent key group* variables. In Figure 3-(ii), there are three connected components and thus three equivalent key group variables $V_1, \dots, V_3$. V_1 represents $A.id$ and $B.Aid$ in this group since Q_2 contains the join condition $A.id = B.Aid$.

The exact computation required to compute the cardinality of the SQL query in Figure 3 extends Equation 1 to sum over the domain $\mathcal{D}$ of each join key.

$$\begin{aligned}
 |Q_2| = & \sum_{a1 \in \mathcal{D}(A.id)} \sum_{a2 \in \mathcal{D}(A.id2)} \sum_{c \in \mathcal{D}(C.id)} \\
 & P_A(A.id = a1, A.id2 = a2 | Q_2(A)) * |Q_2(A)| \\
 & * P_B(B.Aid = a1, B.Cid = c | Q_2(B)) * |Q_2(B)| \\
 & * P_C(C.Aid2 = a2, C.id = c | Q_2(C)) * |Q_2(C)| \\
 & * P_D(D.Cid = c | Q_2(D)) * |Q_2(D)|
 \end{aligned} \tag{2}$$

This computation is extremely expensive — assuming the domains of each id being n , a naive implementation would have the complexity $O(n^3)$. We can improve this complexity to $O(n^2)$ by representing this computation as an equivalent factor graph model [44], which is a particular class of probabilistic graphical models (PGMs) [34]. We have 2 in the exponent because the maximum number of join keys in Q_2 is two in the table A . This is still very expensive, but in Section 3.2, we describe how we can use principled approximations made possible by the factor graph formulation to make the inference complexity practical.

A factor graph is a bipartite graph with two types of nodes: variable nodes, and factor nodes representing an unnormalized probability distribution w.r.t. the variables connected to it. Figure 3-(iii) shows the constructed factor graph $\mathcal{F}$ for computing the cardinality of Q_2 . Specifically, $\mathcal{F}$ contains a variable node for each equivalent key group variable V_i , and a factor node for each table

touched by Q_2 . A factor node is connected to a variable node if the variable represents a key in the table. In this case, the factor node representing table A is connected to V_1 (equivalent to $A.id$) and V_2 (equivalent to $A.id2$). Each factor node maintains an unnormalized probability distribution for the variable nodes connected to it, for e.g., node A maintains $P_A(V_1, V_2 | (Q_2(A))) * |Q_2(A)|$, which is the same as the distribution $P_A(A.id, A.id2 | (Q_2(A))) * |Q_2(A)|$ used in Equation 2. The factor graph model utilizes the graph structure to compute the sum using the well-studied inference algorithms.

We can more formally state this relationship between computing the join cardinalities and factor graphs as the following lemma, whose proof is provided in Section A.2 the technical report [70].


LEMMA 1. *Given a join graph $\mathcal{G}$ representing a query Q , there exists a factor graph $\mathcal{F}$ such that the variable nodes in $\mathcal{F}$ are the equivalent key group variables of $\mathcal{G}$ and each factor node represents a table I touched by Q . A factor node is connected to a variable node if and only if this variable represents a join-key in table I . The potential function of a factor node is defined as table I 's probability distribution of the connected variables (join keys) conditioned on the filter predicates $Q(I)$. Then, calculating the cardinality of Q is equivalent to computing the partition function of $\mathcal{F}$.*

3.2 Inference complexity analysis

Inference on factor graphs is a well-studied problem in the domain of PGMs [34, 44, 46]. Popular approaches to solving this problem are variable elimination (VE) [34] and belief propagation algorithms [35]. In *FactorJoin*'s implementation, we use the VE, which first determines an optimal order of variables and then sums over the distributions defined by the factor graph model in this order. The complexity of VE is $O(N * |D|^{\max(|JK|)})$, where N is the number of equivalent key groups (3 in Figure 3), $|D|$ is the largest domain size of all join keys, and $\max(|JK|)$ is the maximum number of join keys in a single table (2 in Figure 3). Intuitively, this complexity is because the factor nodes need to understand the joint distribution of all join keys in one table. Thus, the largest factor node in the factor graph model maintains a probability distribution of size $|D|^{\max(|JK|)}$. This complexity is not practical for real-world queries, as $|D|$ can be millions, and $\max(|JK|)$ can be larger than 4 in real-world DB instances, such as IMDB [38]. Therefore, instead of calculating the exact cardinality, we propose a new PGM inference algorithm for *FactorJoin* that can estimate an upper bound on cardinality. Past works [1, 3, 5] have shown that cardinality upper bounds can help avoid very expensive query plans since under-estimating a large result can sometimes be catastrophic. For instance, an optimizer might choose to do a nested loop join if a CardEst method under-estimates that the result would easily fit in memory. This would lead to a disastrous plan if the actual result size is much larger than estimated.

Our approximate inference algorithm performs two approximations: 1) reducing the domain size $|D|$ using binning and 2) decreasing the exponent, $\max(|JK|)$, by approximating the distribution of attributes in a single table. Specifically, for binning, we first partition the domain of all join keys into k bins. Instead of summing over the entire domain, we only need to sum over k summarized values (probabilistic bounds) of each bin. The details are provided in Section 4. Second, we approximate the causal relation among all attributes within a table as a tree structure. This allows us to factorize the $\max(|JK|)$ -dimensional joint distribution of all join keys in a table as a product of two-dimensional conditional distributions. This procedure is known as structure learning in Bayesian networks in the PGM domain [6]. Past work on single table cardinality estimation has shown that such tree approximations do not decrease the modeling accuracy significantly [71]. Further details are provided in Section 5.1.

Thereafter, we can directly run the VE inference algorithm on the binned domains and factorized two-dimensional distributions to estimate the cardinality upper bound of Q . With these two

Fig. 4. Workflow of *FactorJoin* framework.

approximate inference techniques, the complexity of estimating the cardinality bound is reduced to $O(N * k^2)$, which is very efficient as both N and k are very small in practice.

3.3 Workflow overview

The workflow of *FactorJoin* contains two phases: offline training and online inference (Figure 4). During the offline phase, *FactorJoin* analyzes the DB instance, bins the domain of join keys, explores the causal relation, and builds CardEst models for every single table. When a join query Q comes in during the online phase, *FactorJoin* formulates the estimation of this query as a factor graph involving only single-table distributions, estimates them using the trained single-table estimators, runs PGM inference, and generates a probabilistic bound for Q . We provide the details as follows.

Offline training phase: Given a new DB instance, *FactorJoin* first analyzes its DB schema and data tables to get all possible join relations between different join-keys. We consider two join-keys to be semantically equivalent if there exists a join relation between them. After identifying all groups of equivalent join-keys, we bin the domains of all join-keys in each group. *FactorJoin* can also optionally use query workload information to optimize this binning procedure (details provided in Section 4.2). Based on the data tables and the binned domains, we explore the causal relation of these attributes and model them as a tree structure (details are provided in Section 5.1). Then, we build the CardEst models to understand the data distribution of every single table. In principle, any single-table CardEst method that is able to provide conditional distributions can be adapted into *FactorJoin*. In practice, we implement two methods: traditional random sampling and the learned data-driven BayesCard method [71]. The sampling method is extremely flexible to use and can easily support any complex filter predicates with disjunctions, string pattern matching, or any user-defined functions. Alternatively, similar to most data-driven methods, BayesCard can only work with queries filtering numeric and categorical attributes, but BayesCard has been shown to consistently produce accurate, fast, and robust estimations for various data distributions [71]. Users of *FactorJoin* can switch between different single-table estimators based on their objectives and knowledge about the DB instances.

Online inference phase: When a join query Q comes in during run-time, *FactorJoin* will first parse its join graph and construct the corresponding factor graph $\mathcal{F}$ (as in Figure 2). Then, it uses the trained single-table estimators to estimate the join-key distributions conditioned on the filter predicates of Q ; *FactorJoin* will use these distributions in the corresponding factor nodes of $\mathcal{F}$. Finally, we run the approximate version of VE inference algorithms on $\mathcal{F}$ to derive the

estimated/probabilistic cardinality bound. We note that the CardEst methods working inside a query optimizer will estimate all sub-plan queries of a target query to decide the best query plan. Since these sub-plan queries contain a large amount of overlapping information, we provide a progressive estimation algorithm to avoid computation redundancy in Section 5.2.

4 PROBABILISTIC BOUND ALGORITHM

In this section, we design a new probabilistic bound algorithm based on binning to calculate the factor graph inference problem. Specifically, we first explain the main ideas with a two-table join example and formulate the algorithm details within *FactorJoin* in Section 4.1. Then, we discuss how to optimize the bin selection using data and query information in Section 4.2.

4.1 Algorithm details

As described in Section 3, the main objective of our probabilistic bound-based algorithm is to reduce the domain size of the join keys, i.e., reducing the complexity of summation in Equation 1. Thus, instead of running the inference algorithm over the entire domain, *FactorJoin* only needs to sum over the binned domain.

Two-table join query example: We use the previous example query Q from Section 3.1 to illustrate the core idea of our probabilistic upper bound algorithm based on binning. The objective is to estimate an upper bound on Q , whose exact computation requires summing over the entire domain of $A.id$ according to Equation 1. Assume that we have a set of k bins $\{bin_1, \dots, bin_k\}$ partitioning the domain of $A.id$ and we apply this same set of bins to $B.Aid$. We need to make sure that a given value in the domain of $A.id$ and $B.Aid$ will always belong to the bin with the same indexes. Then, Equation 1 can be equivalently written as Equation 3. We can then replace the sum over all values in a bin bin_i with a probabilistic bound on bin_i to estimate an upper bound for Q as Equation 4.

$$|Q| = \sum_{i=1}^k \sum_{v \in bin_i} P_A(A.id = v|Q(A)) * |Q(A)| * P_B(B.Aid = v|Q(B)) * |Q(B)| \quad (3)$$

$$\lesssim \sum_{i=1}^k Probabilistic_bound(A, B, bin_i) \quad (4)$$

Motivated by a bound based on the most frequent value (MFV) [22], we use a simple probabilistic bound method to derive $Probabilistic_bound(A, B, bin_i)$ for a particular bin. Specifically, assuming that value $\{a, b, c, e, f\}$ of $A.id$ and $B.Aid$ are binned into bin_1 as shown in Figure 5. We know the summation of all values in bin_1 equals to $8 \times 6 + 4 \times 5 + 3 \times 5 = 83$. This summation has a dominating term of 8×6 , because the count of MFV of bin_1 is 8 for $A.id$ (denoted as $V_1^*(A.id)$) and 6 for $B.Aid$ (denoted as $V_1^*(B.Aid)$) so each value can appear at most 8×6 times in the denormalized table after the join. Since we know the total count of values in bin_1 for $A.id$ is 16, there can be at most $16/8 = 2$ MFVs. Similarly, there can be at most 4 MFVs in bin_1 for $B.Aid$. Therefore, we have the summation of all values in bin_1 is upper bounded by $\min(2, 4) \times 8 \times 6 = 96$.

We formally represent the aforementioned procedure in Equation 5, where $V_i^*(A.id)$ and $P_A(A.id \in bin_i|Q(A)) * |Q(A)|$ are the MFV count and estimated total count of bin_i for $A.id$. Our bound is probabilistic because $P_A(A.id \in bin_i|Q(A)) * |Q(A)|$ is estimated with a single table CardEst method, which may have some errors.




Fig. 5. The bound-based algorithm; each bin is summarized by the most frequent value count and total count.

$$|Q| \lesssim \sum_{i=1}^k \min \left(\frac{P_A(A.id \in bin_i | Q(A)) * |Q(A)|}{V_i^*(A.id)}, \frac{P_B(B.Aid \in bin_i | Q(B)) * |Q(B)|}{V_i^*(B.Aid)} \right) * V_i^*(A.id) * V_i^*(B.Aid) \quad (5)$$

Working with PGM inference: Next, we will explain how to generalize this probabilistic bound to make it work inside the factor graph VE inference algorithm mentioned in Section 3.1. During the offline training, *FactorJoin* computes the MFV counts V^* for all bins of each join-key. During the online phase, *FactorJoin* estimates the probability distribution for the binned domain of all join-keys in each table I . Then, *FactorJoin* puts I 's distribution P_I and MFV counts $V^*(I)$ into the corresponding factor node.

Recall that the VE algorithm generates an optimal elimination order of all variables and sums over the domain of these variables sequentially in this order. Since at each step, the variable elimination algorithm sums out only one variable, its probabilistic bound can be derived using Equation 5 over multiple equivalent join keys. Thus, this algorithm does not need to sum over the entire value domain of a variable but can only sum over the probabilistic bounds in the binned domain instead, which greatly reduces the complexity.

Analysis: This probabilistic bound algorithm significantly reduces the domain size of join keys and enables practical approximate PGM inference. Further, this algorithm can tightly upper-bound the cardinality for most of the time even with a small number of bins, such as $k = 100$. Therefore, as we empirically verify in our evaluation, the estimated cardinality bound is very effective.

4.2 Bin selection optimizations

In our probabilistic bound algorithm, different bins can result in drastically different bound tightness. There are two decisions: how many bins to use for each join key group, and what values to put in the same bin.

Deciding k based on query workloads: The number of bins, k , has a significant effect on the performance of *FactorJoin*: fewer bins aggregate more distinct values in the join key domain to each bin and are thus less accurate but more efficient. An approach to tune k is to set different values of k_i for different equivalent key groups Gr_i . Suppose that users provide a budget K with $K = \sum k_i$. If *FactorJoin* has access to the query workload, it can analyze the join patterns and count the number of times n_i each Gr_i appears in the workload. Our framework can adaptively set a larger k_i for the frequently visited Gr_i and vice versa. In our implementation, we use a simple heuristic — set $k_i = K * n_i / \sum_j n_j$. In this way, we can optimize the modeling capacity of *FactorJoin* to make sure that it is spent on the important joins.

Algorithm 1 Greedy Bin Selection Algorithm (GBSA)

Input: Equivalent key groups $Gr_1, \dots, Gr_m$, where $Gr_i = \{Id_i^1, \dots, Id_i^{|Gr_i|}\}$;
 Column data $\mathcal{D}(Id_i^j)$ of all join keys in the DB instance $\mathcal{D}$;
 Number of bins k_i for each group Gr_i .

```

1: for  $Gr_i \in \{Gr_1, \dots, Gr_m\}$  do
2:    $Bin(Gr_i) \leftarrow []$ 
3:    $Gr'_i \leftarrow \text{sort\_key\_based\_on\_domain\_size}(\mathcal{D}, Gr_i)$ 
4:    $Bin(Gr_i) \leftarrow \text{get\_min\_variance\_bins}(\mathcal{D}(Gr'_i[1]), k_i/2)$ 
5:    $\text{remain\_bins} \leftarrow k_i/2$ 
6:   for  $j \in \{2, \dots, |Gr'_i|\}$  do
7:      $\text{binned\_data} \leftarrow \text{apply\_bin\_to\_data}(\mathcal{D}(Gr'_i[j]), Bin(Gr_i))$ 
8:      $\text{bin\_variance} \leftarrow \text{calculate\_variance}(\text{binned\_data})$ 
9:      $\text{arg\_sort\_idx} \leftarrow \text{arg\_sort\_decreasing}(\text{bin\_variance})$ 
10:    for  $p \in \text{arg\_sort\_idx}[1 : \text{remain\_bins}/2]$  do
11:       $Bin(Gr_i) \leftarrow \text{min\_variance\_dichotomy}(Bin(Gr_i)[p], \text{binned\_data}[p])$ 
12:    end for
13:     $\text{remain\_bins} \leftarrow \text{remain\_bins}/2$ 
14:  end for
15: end for
16: return  $\{Bin(Gr_i) | i, \dots, m\}$ 

```

Binning strategy: For the probabilistic bound algorithm described in Section 4, we observe that the upper bound on a particular bin bin_i can be very loose if the MFV count V_i^* is a large outlier in bin_i . Taking the two table join query as an example, if bin_i contains only one value that appears 100 times in $A.id$ but 10,000 values that only appear once in $B.Aid$, then the bound could be 100 times larger than the actual cardinality. Common binning strategies used in DBMS histograms such as equal-depth or equal-width bins can be catastrophic in such cases. Instead, we want the variance of the join key counts in each bin to be low. This strategy has also been proposed in different contexts [28, 30], but the key challenge is that the same bins will be built for all join keys within an equivalent key group, i.e. join keys that share the same semantics. For example, for a primary key (e.g. “title.id” in IMDB), any binning strategy will have zero variance because of uniqueness. However, the same bins will be applied for its equivalent foreign keys (e.g. “movie_id”) on other tables, which may lead to a large variance depending on how often each value repeats in the other tables. We design a greedy binning selection algorithm (GBSA) to optimize for bins with low variance across all tables, as illustrated in Algorithm 1. Jointly minimizing the variance of one bin for all join keys has exponential complexity. Therefore, GBSA uses a greedy algorithm to iteratively minimize the bin variance for all join keys. At a high level, GBSA first optimizes the minimal variance bins with half budget $k/2$ on the domain of one join key (lines 2-4). Then, it recursively updates these bins by minimizing the variance of other join keys using the other half of the budget (lines 5-13). We put the details of GBSA in Section B our technical report [70].

Discussion: As shown in Section 6.4, GBSA has significantly better performance than naive binning strategies. In the extreme case, if the value counts have zero variance for all equivalent join keys, then our bound can output the exact cardinality. Admittedly, this binning strategy does have the drawback that after applying the filter predicates, the join key frequencies will change and their variance may become higher within a bin. However, we empirically evaluate that for real-world datasets, this phenomenon does not have a severe impact on *FactorJoin*. In future work, we will explore an enhanced binning strategy that can efficiently address this issue.

4.3 Incremental model updates

FactorJoin is friendly for DBs with frequent data changes because we only need to incrementally update the single table statistics. In the following, we will discuss the algorithm details in the scenario of data insertion and data deletion can be handled similarly.

Consider we have trained a stale *FactorJoin* model $\mathcal{F}$ on data $\mathcal{D}$ and would like to incrementally update $\mathcal{F}$ with the inserted data $\mathcal{D}'$. *FactorJoin* needs to identify the inserted tuple values from $\mathcal{D}'$ for every join key and put them into the original bins optimized on $\mathcal{D}$. Then, it will update the total and most-frequent value count in each bin. At last, *FactorJoin* can incrementally update the base table CardEst models with off-the-shelf tools (e.g. *BayesCard* provides an efficient and effective approach to incrementally update the single-table models using $\mathcal{D}'$ [71]; materializing a new sample on $\mathcal{D} \cup \mathcal{D}'$ incrementally updates the sampling-based estimator).

Therefore, the incremental update of *FactorJoin* does not need to normalize the joined tables nor need to use executed queries. It is also empirically verified to be very efficient and effective. However, during incremental updates, the *FactorJoin* models keep the same set of bins, which is optimized on the previous data and thus might not be optimal after data insertion. This might cause a slight degradation in performance and the user can choose to retrain the model, which is also not very costly, in case of massive data updates.

5 IMPROVING FACTORJOIN EFFICIENCY

We describe two techniques to improve the modeling and estimation efficiency of *FactorJoin*. First, we model the joint distribution of attributes on a single table as a tree structure and use it to simplify the PGM inference calculations. Next, we define a systematic way of reusing estimates of intermediate sub-plan queries to estimate a larger query. This is only possible because we decompose cardinality estimates into single table sub-components.

5.1 Exploring attributes causal relation

Recall from Section 3.2, the factor graph computation to calculate the cardinality bound is $O(N * k^{\max(|JK|)})$, where N is the number of tables, k is the number of bins, and $\max(|JK|)$ is the maximum number of join keys in a table. The exponent term, $\max(|JK|)$, represents the factor nodes requiring the joint distribution between all join keys on a single table. As $\max(|JK|)$ can sometimes be large (e.g., in IMDB it can be up to 4), we need to reduce the dimensionality of the distributions defined in factor nodes to ensure an efficient CardEst procedure. Fortunately, real-world data has a lot of correlations that can be used to simplify this joint distribution.

Consider a table A with 6 attributes, $\{id_1, id_2, id_3, id_4, attr_1, attr_2\}$, four of which are join keys. The factor graph computation may need to estimate quantities such as:

$$P_A(id_1, id_2, id_3, id_4 | Q(A)) \quad (6)$$

This requires storing a 4 dimensional probability distribution, and even with binning the domains of all ids into k bins, it still requires k^4 space and inference time complexity. This joint distribution can be represented as a graph where each attribute is a node, and there are edges between every pair of attributes. We will use the dependencies and relationships between the attributes and model them as a Bayesian network (BN) [34]. Specifically, we will assign each edge in the joint distribution graph a weight proportional to the mutual information between the two attributes connected by the edge. Then, we will remove edges with the least mutual information until only a tree structure graph remains to represent the joint probability over these attributes. This process is known as the

Chow-Liu algorithm for decomposing joint distributions [6]. The factorized tree representation of the joint distribution can let us approximate Equation 6 with an equation of the form:

$$P_A(id_1|Q(A)) * P_A(id_2|Q(A)) * P_A(id_3|id_1) * P_A(id_4|id_3) \quad (7)$$

, which is much more efficient to compute because it only involves a product over at most a two-dimensional distribution. This dimension reduction can cause some errors. But it has been shown empirically that the factorized distributions are accurate approximations of the original distribution for most real-world datasets [71].

Using such an optimization, instead of storing the $\max(|JK|)$ -dimensional distributions in the factor nodes, the factor graph only needs to store a series of one or two-dimensional data distributions without a significant decrease in estimation accuracy. Therefore, the overall complexity for running approximate inference on this factor graph becomes $O(N * k^2)$ in both storage and inference speed. Since k is relatively small (around one hundred), estimating one query using *FactorJoin* is very efficient.

In principle, *FactorJoin* can also use other methods for the single table distribution inference. This may involve traditional methods, such as sampling, or more powerful graph-structured factorizations of the distributions — which can increase modeling accuracy at the cost of slower estimation speed. Thus, users of *FactorJoin* can make trade-offs between estimation accuracy and efficiency. We leave exploring these trade-offs as future work.

5.2 Progressive estimates of sub-plan queries

In order to optimize the plan for one query, the CardEst method needs to estimate the cardinalities of hundreds or thousands of sub-plan queries. Estimating all these queries independently is very inefficient because they contain a lot of redundant computation. Similar to the traditional CardEst methods deployed in real DBMSes, *FactorJoin* supports estimating all sub-plan queries of one query progressively, reusing the computation as much as possible. At a high level, the progressive estimation algorithm will estimate all sub-plan queries in a bottom-up fashion. *FactorJoin* first estimates all two-table join sub-plan queries, then it estimates the three-table join sub-plan queries, which will contain the two-table joins as their sub-structures, and so on.

Joining Factor Graphs: The key insight is that when we estimate the cardinality of a two-table join, we can also combine their factor graphs into a new factor graph for the joined table. Specifically, as described in Section 3, each base table is represented as a factor node in the factor graph. For instance, the factor node representing A will store the unnormalized probability distribution and the MFV counts for the binned domain of join keys in A . When estimating the cardinality of joining A and C , our approximate inference algorithm in Equation 5 computes the bound over all bins and sums them. These bounds essentially define an unnormalized probability distribution over all join keys on the denormalized table $A \bowtie C$. We can derive a bound on the new MFV counts for a join-key $A \bowtie C$ as the multiplication of the corresponding MFV counts in A and C . Therefore, we can cache this new probability distribution and MFV counts in a new factor node representing $A \bowtie C$. Thus, when estimating the cardinality of $A \bowtie B \bowtie C$ we can consider it as a join of two single tables ($A \bowtie C$) and B .

Since all sub-plans can essentially be considered a join of two sub-plans that we have already estimated, our inference algorithm only needs to compute the join of two-factor nodes for each sub-plan query. Therefore, the progressive estimation algorithm has no redundant computation and maximizes efficiency. In experiments, with this algorithm, *FactorJoin* can estimate the cardinality of 10,000 sub-plan queries in one second, which is more than ten times faster than estimating all these queries independently. Note that the progressive estimation algorithm is only possible to implement for *FactorJoin* and the traditional methods because it requires decomposing join queries

into single-table estimates. None of the existing learned methods can apply this method since their models are built to model the distributions of specific join patterns.

6 EXPERIMENTAL EVALUATION

In this section, we empirically demonstrate the advantages of our framework over existing CardEst methods. We first introduce the datasets, baselines, and experimental environment in Section 6.1. Then, we explore the following questions:

- **Overall performance** (Section 6.2): how much improvement can *FactorJoin* achieve in terms of end-to-end query time? What are the model size and training time compared to SOTA methods?
- **Detailed analysis** (Section 6.3): Why does *FactorJoin* attain good performance? When does it perform better or worse than existing baselines? How will it perform in terms of data updates?
- **Ablation study** (Section 6.4): How much does each optimization technique of *FactorJoin* contribute to the overall performance?

6.1 Experimental Setup

Datasets: We evaluate the performance of our CardEst framework on two well-established benchmarks: *STATS-CEB* [18] and *IMDB-JOB* [38], whose statistics are summarized in Table 2.

The *STATS-CEB* benchmark uses the real-world dataset STATS, which is an anonymized dump of user-contributed content on Stack Exchange (Stack Overflow and related websites). This dataset consists of 8 tables with 34 columns and more than one million rows. It contains a large number of attribute correlations and skewed data distributions. The *STATS-CEB* query workload contains 146 queries involving 70 different join templates with real-world semantics. The true cardinalities of these queries range from 200 to $2 \cdot 10^{10}$, with some queries taking more than one hour to execute on Postgres. The *STATS-CEB* benchmark only has star or chain join templates and numerical or categorical filtered attributes — therefore it is possible to evaluate all existing CardEst baselines.

The *IMDB-JOB* benchmark uses the real-world Internet movie database (IMDB). This dataset contains 21 tables with more than fifty million rows. The *IMDB-JOB* query workload contains 113 queries with 33 different join templates. Some queries have more than 10,000 sub-plan queries, posing a significant challenge to the efficiency of CardEst methods. Because *IMDB-JOB* contains cyclic joins and string pattern matching filters, the *JoinHist* and existing learned data-driven methods do not support this benchmark.

The queries in these two benchmarks cover a wide range of joins, filter predicates, cardinalities, and runtimes. To the best of our knowledge, they represent the most comprehensive and challenging real-world benchmarks for CardEst evaluation.

Baselines: We compare our framework with the most competitive baselines in each CardEst method category.

- 1) *PostgreSQL* [10] refers to the histogram-based CardEst method used in PostgreSQL.
- 2) *JoinHist* [7] is the classical join-histogram method whose details are explained in Section 2.2.
- 3) *WJSample* [40] uses a random walk based-method called wander join to generate samples from the join of multiple tables. Its superior performance over traditional sampling-based methods has been demonstrated in recent studies [40, 56, 77]. We limit the amount of time spent for each random walk of *WJSample* so the overall time spent for estimating cardinalities is comparable to other methods. If we allow *WJSample* longer inference time, it can generate more effective query plans — but it is impractical since its total estimation time will exceed the query execution time.
- 4) *MSCN* [33] is a learned query-driven method based on a multi-set convolutional network model. On both datasets, we train MSCN on roughly 100K sub-plan training queries that have similar distributions to the testing query workloads.

Category	Statistics	STATS-CEB	IMDB-JOB
Data	# of tables	8	21
	# of rows in each table	$10^3 - 4 \cdot 10^5$	$6 - 3 \cdot 10^7$
	# of columns in each table	3 – 10	2-12
	# of join keys	13	36
	# of equivalent key groups	2	11
Query	# of queries	146	113
	# of join templates	70	33
	join template type	star & chain	+cyclic
	# of filter predicates	1 – 16	1 – 13
	filter attributes	numerical & categorical	+string LIKE
	# of sub-plan queries	1 – 75	$8 - 1 \cdot 10^4$
	true cardinality range	$200 - 2 \cdot 10^{10}$	$1 - 4 \cdot 10^6$
	Postgres runtime (s)	$0.05 - 10^4$	$5 \cdot 10^{-4} - 450$

Table 2. Summary of STATS-CEB and IMDB-JOB benchmark.

5) *BayesCard* [71], 6) *DeepDB* [25], and 7) *FLAT* [78] are the learned data-driven CardEst methods. They use fanout methods to understand the data distributions of all join templates. Then, they apply different distribution models (Bayesian networks, sum-product networks [57], factorized sum-product networks [73], respectively) to capture the data distribution on single tables or denormalized join tables. We test these methods on *STATS-CEB* using the optimally-tuned model parameters, as used in a prior benchmark study [18].

8) *PessEst* [5] is the SOTA bound-based method, which leverages randomized hashing and data sketching to tighten the bound for join queries. Its estimation has been verified to be very effective in real-world DBMSes. We use the best-tuned hyperparameters to evaluate *PessEst* on both benchmarks.

9) *U-Block* [22] uses top-k statistics to estimate cardinality bounds. The original paper designed a new plan enumeration strategy to work with this bound and showed promising results on *IMDB-JOB*. However, for a fair comparison, we only evaluate the standalone *U-Block* method.

10) *TrueCard* outputs true cardinality for given queries with no estimation latency. This baseline does not refer to any specific method but represents the optimal CardEst performance.

We omit comparisons with many other CardEst methods [4, 11, 12, 20, 21, 31, 32, 39, 61, 62, 68, 69, 74] because the baselines above have demonstrated clear advantages over these methods in various aspects [18]. In addition, some newly proposed methods [42, 63, 67] do not have open-source implementations yet.

We implemented *FactorJoin* in Python mainly using the NumPy package. Regarding the hyperparameters of *FactorJoin*, we assume no knowledge of the query workload and set the bin size k to 100 for both datasets. We use the BN-based CardEst method [71] as the single-table estimators for *STATS-CEB* benchmark because this method has been shown to be very accurate, fast, and robust for single-table estimations. For *IMDB-JOB*, we use random sampling with 1% sampling rate as the single-table estimator because this benchmark contains very complicated filter predicates such as disjunctions and string pattern matching, which are not supported by the learned data-driven methods. We discuss and compare other hyperparameters of our framework in Section 6.4.

Environment: We use the open-source implementation of all the baselines and tune their parameters to achieve the best possible performance. All experiments are performed on a Ubuntu server with Intel Core Processor CPU with 20 cores, 57GB DDR4 main memory, and 300GB SSD. For end-to-end evaluation, we integrate our framework along with all the baselines into the query




Fig. 6. Overall performance on STATS-CEB and IMDB-JOB.

optimizer of PostgreSQL 13.1 following the procedure described in the recent work [17, 18]. We build the primary key and foreign key indices and disable parallel execution in PostgreSQL. Specifically, we inject into PostgreSQL all sub-plan query cardinalities estimated by each method, so the PostgreSQL optimizer uses the injected cardinalities to optimize the query plan, and record the end-to-end query runtime (planning plus execution time). To eliminate the effects of cache, we first execute each workload multiple times. Then, we execute the queries three times for each method and report the average and the standard errors.

6.2 Overall performance

The most straightforward criterion to evaluate the performance of a CardEst method is to test how much each method improves the end-to-end query performance inside the query optimizer. Therefore, we compare the end-to-end performance along with the model size and training time of existing baselines on *STATS-CEB* and *IMDB-JOB*. We summarize the overall performance of all baselines in Figure 6. Recall that *JoinHist*, *BayesCard*, *DeepDB*, and *FLAT* can not support the *IMDB-JOB* benchmark. In addition, we do not provide the model size and training time for *Postgres*, *TrueCard*, *JoinHist*, *WJSample*, and *U-Block* as they are negligible.

Performance on STATS-CEB: In addition to Figure 6, we provide a detailed comparison of query execution and planning time (sum over all queries in the *STATS-CEB*) in Table 3, with the relative end-to-end improvement over *Postgres* shown in the last column, i.e. $(\text{Postgres time} - \text{method time}) / \text{Postgres time}$.

FactorJoin achieves the best end-to-end query runtime (19, 116s), which is close to optimal performance (18, 456s for *TrueCard*). When compared with the traditional methods (*Postgres*, *JoinHist*, and *WJSample*), *FactorJoin* has significantly better query execution time, indicating that our estimates are much more effective at generating high-quality query plans. Meanwhile, the planning time of our framework is close to *Postgres*.

Judging from the pure query execution time in Table 3, *FactorJoin* has comparable estimation effectiveness as the previous SOTA method *FLAT* on this benchmark. This learned data-driven

Method	End-to-end time	Exec. + plan time	Improvement
Postgres	35,341s	35,316s + 25s	–
TrueCard (Optimal)	18,456s	18,432s + 24s	47.8%
JoinHist	33,201s	33,173s + 28s	6.1%
WJSample	59,481s	59,436s + 45s	-68.4%
MSCN	25,559s	25,524s + 35s	27.7%
BayesCard	22,679s	22,644s + 35s	35.9%
DeepDB	20,472s	20,304s + 168s	42.0%
FLAT	19,331s	18,934s + 437s	45.3%
PessEst	21,006s	19,872s + 1,135s	40.5%
U-Block	38,618s	38,592s + 26s	-9.3%
FactorJoin (Ours)	19,116s	19,080s + 36s	45.9%

Table 3. End-to-end performance on STATS-CEB.

Method	End-to-end time	Exec. + plan time	Improvement
Postgres	2,472s	2,451s + 21s	–
TrueCard (Optimal)	375s	358s + 17s	84.8%
WJSample	13,618s	12,445s + 1,173s	-450.9%
MSCN	2,039s	1,937s + 102s	18.1%
PessEst	4,044s	1,230s + 2,814s	-63.6%
U-Block	2,792s	2,762s + 30s	-12.9%
FactorJoin (Ours)	1,324s	1,272s + 42s	46.4%

Table 4. End-to-end performance on IMDB-JOB.

method uses the fanout method to understand the distribution of all join patterns, which is accurate but very inefficient. Specifically, *FactorJoin* can achieve better end-to-end performance than *FLAT*, with $10\times$ smaller planning time, $160\times$ smaller model size, and $240\times$ faster training time. When compared to the other learned methods: *BayesCard*, *DeepDB*, and *MSCN*, *FactorJoin* can simultaneously achieve better query execution and planning time, smaller model size, and faster training time. Therefore, we empirically verify that instead of modeling the distributions of all join patterns, accurately decomposing the join query into single-table estimates can generate equally effective but more efficient estimation with much smaller model size and training time.

FactorJoin and the SOTA bound-based method *PessEst* can both generate very effective estimates. This observation verifies that a bound-based method can produce effective estimates because it can help avoid very expensive query plans [1, 3, 5, 18]. However, the planning time of *PessEst* (1135s) is undesirable because it needs to materialize tables after applying the filters and populate the upper bound during run-time. *FactorJoin* has $35\times$ less planning latency since instead of using exact bounds, we use the much more efficient probabilistic bounds derived from single-table estimates.

Performance on IMDB-JOB: Table 4 shows a detailed end-to-end time comparison for the *IMDB-JOB* workload. *FactorJoin* achieves the best end-to-end query runtime (1324s) amongst all baselines. Similar observations hold as in the *STATS-CEB* workload when compared to the traditional methods. Like *STATS-CEB*, *PessEst* achieves the best performance in terms of pure execution time but has a much larger overhead and planning latency. *FactorJoin* has comparable performance as *PessEst* in terms of pure execution time. However, our planning time is $50\times$ shorter than *PessEst*, making our framework much faster in overall end-to-end query time. Specifically, our framework can estimate 10,000 sub-plan queries within one second, which is close to *Postgres*.




Fig. 7. Relative estimation errors on *STATS-CEB*.

However, unlike in *STATS-CEB* where we have near-optimal performance, there still exists a large gap between *FactorJoin* (1324s) and the optimal *TrueCard* (375s). There are two possible explanations. First, our probabilistic bound error can accumulate with the number of joins. Compared to *STATS-CEB*, the *IMDB-JOB* contains more than twice the number of joins in the queries. Therefore, the relative performance of our framework drops severely when shifting from *STATS-CEB* to *IMDB-JOB*. Second, the base-table filters in *IMDB-JOB* contain a large amount of highly-selective predicates, making our sampling-based single-table CardEst less accurate. These observations suggest room for improving *FactorJoin*, which we discuss further in the future work section.

Summary: *FactorJoin* achieves the best performance among all the baselines on both benchmarks. Specifically, *FactorJoin* is as effective as the previous SOTA methods, and simultaneously as efficient and practical as the traditional CardEst methods.

6.3 Detailed Analysis


In this section, we first study why *FactorJoin* is able to achieve SOTA performance by analyzing the tightness of our probabilistic bound. Then, we exhaustively compare the CardEst method's end-to-end performance on queries with different runtimes to investigate when can *FactorJoin* perform better or worse than the competitive baselines. For compactness, here, we only report the results of the most representative methods (*Postgres*, *FLAT*, *PessEst*) on *STATS-CEB*; the remaining results are in Section C of our technical report [70].

Bound tightness: We show the relative errors between the baselines' estimates and the true cardinalities (estimate / true) for all sub-plan queries on *STATS-CEB* in Figure 7. Overall, all three SOTA methods significantly outperform *Postgres* in terms of estimation accuracy. *PessEst* generates exact upper bounds and never underestimates. *FLAT* uses a much larger model to understand the distributions of join patterns and produces the most accurate estimates.

FactorJoin can output an upper bound on cardinality for more than 90% of the sub-plan queries. Most of the marginal underestimates are very close to the true cardinality, which can still generate relatively effective plans. Our generated bounds are slightly tighter than the bounds from *PessEst*, which explains why our query plans achieve slightly better results.

Detailed comparison: We sort the 146 queries of *STATS-CEB* based on their *Postgres* end-to-end runtime and cluster them into 6 intervals. Figure 8 reports relative improvements of the competitive baselines over *Postgres* for each query on the left and each cluster of queries on the right.

For the very short-running queries (which represent an OLTP-like workload), *Postgres* is the best amongst all baselines. These baselines perform worse because the estimation latency plays a significant role in these queries. We observe that improving the estimation accuracy of these queries has a very limited effect on the query plan quality. This also explains why the optimal

Fig. 8. Per query performance of *STATS-CEB*.

Method	Update time	End-to-end time	Improvement
BayesCard	84s	22,679s	35.9%
DeepDB	310s	21,352s	39.6%
FLAT	422s	22,120s	37.4%
FactorJoin (Ours)	2.5s	20,015s	43.4%

Table 5. Incremental update performance on *STATS-CEB*.

TrueCard only marginally outperforms *Postgres* on queries with less than 1s of runtime. However, these queries only contribute a negligible proportion of total runtime in *STATS-CEB*. The query optimizer should fall back to default traditional CardEst methods for OLTP-like workload. *FLAT* and *PessEst* are significantly worse than *FactorJoin* because their planning latency surpasses the execution time, which overshadows the minor improvement in query plans.

For the extremely long-running queries, the advantages of the SOTA methods over *Postgres* gradually appear. The reason is that estimation latency becomes increasingly insignificant for queries with a long execution time. Thus, *FLAT* and *PessEst* which spend a long time planning can generate significantly better query plans. *FactorJoin* has comparable performance on these queries.

Incremental updates: Following the same evaluation setup as previous work [18], we train a stale *FactorJoin* model on *STATS* data created before 2014 (roughly 50%) and insert the rest of the data to incrementally update this model.

We provide the end-to-end query time of the updated *FactorJoin* and the update time on *STATS-CEB* queries in Table 5. Since *FactorJoin* only needs to update single-table statistics, it is extremely efficient, i.e. taking only 2.5s to update millions of tuples. This update speed up to 168× faster than the learned data-driven methods with better end-to-end query performance after model update. We cannot fairly compare with the learned query-driven methods because there does not exist a training query workload after the data insertion. It worth noticing that the update times reported for other methods in Table 5 include the time to re-compute the denormalized join tables for the inserted data, so numbers are larger than the ones reported in the original paper [18].

However, we do see a slight drop in the end-to-end improvement when compared to the *FactorJoin* trained on the entire dataset (43.4% in Table 5 versus 45.9% in Table 3). This small performance difference is due to the fact that the bins are decided on the data before 2014 and remain fixed during incremental updates. Therefore, after inserting the new data, the min-variance property of the bins could be violated, resulting in less accurate predictions.

6.4 Ablation study

We conduct a series of ablation studies to analyze different optimization techniques inside *FactorJoin*. Specifically, we first investigate how the number of bins (k) affects the overall performance. Next,




Fig. 9. Performance for different number of bins.

we analyze the effectiveness of our new bin selection algorithm (GBSA in Section 4.2) over the traditional equal-width and equal-depth algorithms. Then, we study the performance of plugging-in different single-table CardEst methods into *FactorJoin*. Due to space limitations, we only provide the results on *STATS-CEB*. At last, we investigate how removing each of the simplifying assumptions improves the performance of the original joining histogram.

Number of bins (k): We train and evaluate *FactorJoin* with five different number of bins (1, 10, 50, 100, 200). Figure 9 shows the end-to-end query time (A), bound tightness (B), estimation latency per query (C), training time (D), and model size (E) of *FactorJoin*. For bound tightness, we report the 50%, 95%, 99% percentiles of all queries' relative error: estimated/true, as the one used in Figure 7.

Figure 9-(A-B) shows that larger numbers of bins (k) will generate tighter cardinality bounds and thus more effective query plans. Figure 9-(C-E) shows that more bins also increase estimation latency, training time, and model size, as expected. There are several interesting observations from these results.

- According to Figure 9-(B), with $k = 1$ (i.e. bin the entire domain of each join key into one bin) *FactorJoin* does not generate tight bounds but it still outperforms *Postgres* baseline by 24.9%. This highlights the importance of upper bounds over underestimation and verifies the effectiveness of bound-based algorithms.

- Increasing k can substantially increase the bound tightness. However, at some point (from $k = 100$ to $k = 200$) the framework stops achieving better end-to-end performance. One explanation is that this fine-grained improvement in estimation accuracy can generate a slightly better query plan but the slower estimation latency cancels out the improvements in query plans.

- The model size increases quadratically with k because *FactorJoin* has $O(N * k^2)$ complexity in both storage and inference. However, estimation latency increases linearly with k because the overall inference time is dominated by the single-table CardEst methods' inference time, whose complexity is linear with k .

Different bin selection algorithms: Recall that in Section 4.2, we design a new algorithm (GBSA) to optimize the binning process for *FactorJoin*. Here, we compare it with the traditional equal-width and equal-depth binning algorithms. We set $k = 100$ for all three algorithms. They have approximately the same estimation latency, training time, and model size. Thus, we only report the

Algorithm	End-to-end time	Improvement	Bound Tightness (rel. error)		
			50%	95%	99%
Equal-width	23,868s	32.4%	8.7	3,135	2×10^5
Equal-depth	23,436s	33.6%	8.4	2,050	7×10^4
GBSA	19,116s	45.9%	3.3	44	2,782

Table 6. Performance of different binning techniques.

Method	End-to-end time	Exec. + plan time	Improvement
BayesCard	19,116s	19,080s + 36s	45.9%
Sampling	20,633s	20,592s + 41s	41.6%
TrueScan	19,334s	18,756s + 578s	45.3%

Table 7. *FactorJoin*, varying single-table CardEst methods.

end-to-end performance and bound tightness in Figure 6. We observe that GBSA generates much tighter upper bounds, leading to significantly better end-to-end performance. This demonstrates the effectiveness of our GBSA algorithm.

Varying single-table CardEst methods: To compare the importance of different single-table CardEst methods in *FactorJoin*, we tried three different estimators: 1) Bayesian network (*BayesCard*) for single-table, which is the SOTA learned single-table method [71], 2) *Sampling*, which draws a random sample (5%) from single tables on-the-fly and estimates the cardinalities of filter predicates, and 3) *TrueScan* which scans and filters the tables during query time and calculate the true cardinalities. The end-to-end performance of different single table estimators with bin size $k = 100$ is shown in Table 7. The *BayesCard* method performs significantly better than *Sampling* because of more accurate estimates and slightly faster estimation speeds. *TrueScan* produces an exact upper bound on the cardinalities and thus generates more effective query plans. However, the estimation latency is too high, making its overall end-to-end performance worse than *BayesCard*.

Improvement over joining histograms: Since *FactorJoin* follows the convention of *JoinHist*, we investigate how much each component of *FactorJoin* help improve this method. Specifically, apart from the original *JoinHist* method, we evaluate and compare the following variants. We incorporate our new probabilistic bound into *JoinHist* to remove its *join uniformity assumption* (denote as *with Bound*). We incorporate into *JoinHist* the histograms learned from *BayesNet* that represent conditional distributions of join keys given the filter predicates (denote as *with Conditional*). This variant avoids the *attribute independent assumption*. *FactorJoin* reduces to the *JoinHist* with both bound and conditional techniques on non-cyclic join templates. The results in Table 8 demonstrate that we can achieve significant improvement over the *JoinHist* method by removing either or all of its simplifying assumptions.

Summary: With a reasonable bin size k , *FactorJoin* generates tight bounds and effective query plans. The GBSA bin selection algorithm is very effective at generating high-quality plans. Single-table estimators do have an impact on *FactorJoin*'s overall performance and we should choose one that can generate effective and efficient estimates. Overall, all different settings/hyperparameters of *FactorJoin* can significantly outperform *Postgres* as well as all existing SOTA methods, demonstrating the robustness and stability of *FactorJoin*.

7 RELATED WORK

We briefly review literatures on single-table CardEst methods and the learned query optimizers.

Single-table CardEst: Machine learning approaches have been use to solve single-table CardEst problem, which can achieve accurate estimation with very low latency and overhead [71, 75,

Method	End-to-end time	Exec. + plan time	Improvement
JoinHist	33,201s	33,173s + 28s	6.1%
with Bound	29,175s	29,147s + 28s	17.46%
with Conditional	23,450s	23,414s + 36s	31.7%
with Both (Ours)	19,116s	19,080s + 36s	45.9%

Table 8. Performance comparison with JoinHist.

78]. Specifically, the data-driven learned methods build statistical models to understand data distributions, such as deep autoregressive models [19, 75], Bayesian networks [13, 16, 71], sum-product-networks [25], factorized sum-product-networks [78], and normalizing flow models [67]. As for the query-driven methods, the first approach using neural networks for CardEst has been proposed for UDF predicates [36]. Later on, a semi-automatic alternative [47] and a regression-based model [2] were used for general predicates. Recently, more sophisticated supervised models such as multi-set convolutional networks [33], XG-boost [11], tree-LSTM [63], and deep ensembles [42], are used to provide accurate estimation.

These works are orthogonal to our framework and in principle, we can support plugging in any one of them for single table CardEst.

Learned query optimizer: Apart from CardEst, there exist a large number of works that use ML to solve other tasks of query optimizer, such as cost estimation and join order selection. Learned cost estimation methods use tree convolution [51] and tree-LSTM [63] to encode a query plan as a tree and map the encoding to its estimated costs. Active learning [45] and zero-shot learning [24] are proposed to tackle this problem from a new perspective. Some deep reinforcement learning approaches [50, 76] have been proposed to determine the optimal join order. Recently, many methods [48, 49, 53, 72] propose to learn the entire query optimizer end-to-end without a clear separation of these components.

8 CONCLUSIONS

In this paper, we propose *FactorJoin*, a framework for cardinality estimation of join queries. It combines classical join-histogram methods with learned single-table cardinality estimates into a factor graph model. This framework converts CardEst problem into an inference problem over the factor graph involving only single-table distributions. We further propose several optimizations to make this problem tractable for large join graphs. Our experiments show that *FactorJoin* generates effective and efficient estimates and is suitable for system deployment on large real-world join benchmarks. Specifically, *FactorJoin* produces more effective estimates than the previous SOTA learned methods, with 40x less estimation latency, 100x smaller model size, 100x faster updating speed, and 100x faster training speed, while matching or exceeding them in terms of query execution time. We believe this work points out a new direction for estimating join queries, which would enable truly practical learned CardEst methods as a system component.

9 ACKNOWLEDGMENTS

This research is supported by the MIT Data Systems and AI Lab (DSAIL) and NSF IIS 1900933.

REFERENCES

- [1] Mahmoud Abo Khamis, Hung Q Ngo, and Dan Suciu. 2017. What do Shannon-type Inequalities, Submodular Width, and Disjunctive Datalog have to do with one another?. In *Proceedings of the 36th ACM SIGMOD-SIGACT-SIGAI Symposium on Principles of Database Systems*. 429–444.

- [2] Mert Akdere, Ugur Cetintemel, Matteo Riondato, Eli Upfal, and Stanley B Zdonik. 2012. Learning-based query performance modeling and prediction. In *2012 IEEE 28th International Conference on Data Engineering*. IEEE, 390–401.
- [3] Albert Atserias, Martin Grohe, and Dániel Marx. 2008. Size bounds and query plans for relational joins. In *2008 49th Annual IEEE Symposium on Foundations of Computer Science*. IEEE, 739–748.
- [4] Nicolas Bruno, Surajit Chaudhuri, and Luis Gravano. 2001. STHoles: a multidimensional workload-aware histogram. In *SIGMOD*. 211–222.
- [5] Walter Cai, Magdalena Balazinska, and Dan Suciu. 2019. Pessimistic cardinality estimation: Tighter upper bounds for intermediate join cardinalities. In *SIGMOD*. 18–35.
- [6] C. Chow and Cong Liu. 1968. Approximating discrete probability distributions with dependence trees. *IEEE transactions on Information Theory* 14, 3 (1968), 462–467.
- [7] Alberto Dell’Era. 2007. Join Over Histograms. Available on www.adellera.it/investigations/join_over_histograms (2007).
- [8] Alberto Dell’Era. 2017. Oracle Database Online Documentation 12c Release 1 (12.1). Available at: <http://docs.oracle.com/database/121/index.html> (2017).
- [9] Amol Deshpande, Minos Garofalakis, and Rajeev Rastogi. 2001. Independence is good: Dependency-based histogram synopses for high-dimensional data. *ACM SIGMOD Record* 30, 2 (2001), 199–210.
- [10] PostgreSQL Documentation 12. 2020. Chapter 70.1. Row Estimation Examples. <https://www.postgresql.org/docs/current/row-estimation-examples.html> (2020).
- [11] Anshuman Dutt, Chi Wang, Azade Nazi, Srikanth Kandula, Vivek Narasayya, and Surajit Chaudhuri. 2019. Selectivity estimation for range predicates using lightweight models. *PVLDB* 12, 9 (2019), 1044–1057.
- [12] Dennis Fuchs, Zhen He, and Byung Suk Lee. 2007. Compressed histograms with arbitrary bucket layouts for selectivity estimation. *Information Sciences* 177, 3 (2007), 680–702.
- [13] Lise Getoor, Benjamin Taskar, and Daphne Koller. 2001. Selectivity estimation using probabilistic models. In *SIGMOD*. 461–472.
- [14] Dimitrios Gunopulos, George Kollios, Vassilis J Tsotras, and Carlotta Domeniconi. 2000. Approximating multi-dimensional aggregate range queries over real attributes. In *SIGMOD*. 463–474.
- [15] Dimitrios Gunopulos, George Kollios, Vassilis J Tsotras, and Carlotta Domeniconi. 2005. Selectivity estimators for multidimensional range queries over real attributes. *The VLDB Journal* 14, 2 (2005), 137–154.
- [16] Max Halford, Philippe Saint-Pierre, and Franck Morvan. 2019. An approach based on bayesian networks for query selectivity estimation. *DASFAA* 2 (2019).
- [17] Yuxing Han. 2021. Github repository: E2E benchmark. <https://github.com/Nathaniel-Han/End-to-End-CardEst-Benchmark> (2021).
- [18] Yuxing Han, Ziniu Wu, Peizhi Wu, Rong Zhu, Jingyi Yang, Liang Wei Tan, Kai Zeng, Gao Cong, Yanzhao Qin, Andreas Pfadler, et al. 2021. Cardinality Estimation in DBMS: A Comprehensive Benchmark Evaluation. *VLDB Endowment* (2021).
- [19] Shohedul Hasan, Saravanan Thirumuruganathan, Jeess Augustine, Nick Koudas, and Gautam Das. 2019. Multi-attribute selectivity estimation using deep learning. In *SIGMOD*.
- [20] Shohedul Hasan, Saravanan Thirumuruganathan, Jeess Augustine, Nick Koudas, and Gautam Das. 2020. Deep Learning Models for Selectivity Estimation of Multi-Attribute Queries. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*. 1035–1050.
- [21] Max Heimel, Martin Kiefer, and Volker Markl. 2015. Self-tuning, gpu-accelerated kernel density models for multidimensional selectivity estimation. In *SIGMOD*. 1477–1492.
- [22] Axel Hertzschuch, Claudio Hartmann, Dirk Habich, and Wolfgang Lehner. 2021. Simplicity Done Right for Join Ordering. *CIDR* (2021).
- [23] Benjamin Hilprecht. 2019. Github repository: deepdb public. <https://github.com/DataManagementLab/deepdb-public> (2019).
- [24] Benjamin Hilprecht and Carsten Binnig. 2022. One Model to Rule them All: Towards Zero-Shot Learning for Databases. *CIDR* (2022).
- [25] Benjamin Hilprecht, Andreas Schmidt, Moritz Kulesa, Alejandro Molina, Kristian Kersting, and Carsten Binnig. 2019. DeepDB: learn from data, not from queries!. In *PVLDB*.
- [26] Yannis Ioannidis. 2003. The history of histograms (abridged). In *Proceedings 2003 VLDB Conference*. Elsevier, 19–30.
- [27] Yannis E Ioannidis. 1993. Universality of serial histograms. In *VLDB*, Vol. 93. Citeseer, 256–267.
- [28] Yannis E Ioannidis and Stavros Christodoulakis. 1991. On the propagation of errors in the size of join results. In *Proceedings of the 1991 ACM SIGMOD International Conference on Management of data*. 268–277.
- [29] Yannis E Ioannidis and Stavros Christodoulakis. 1993. Optimal histograms for limiting worst-case error propagation in the size of join results. *ACM Transactions on Database Systems (TODS)* 18, 4 (1993), 709–748.
- [30] Yannis E Ioannidis and Viswanath Poosala. 1995. Balancing histogram optimality and practicality for query result size estimation. *Acm Sigmod Record* 24, 2 (1995), 233–244.

- [31] Andranik Khachatryan, Emmanuel Müller, Christian Stier, and Klemens Böhm. 2015. Improving accuracy and robustness of self-tuning histograms by subspace clustering. *IEEE TKDE* 27, 9 (2015), 2377–2389.
- [32] Martin Kiefer, Max Heimerl, Sebastian Breß, and Volker Markl. 2017. Estimating join selectivities using bandwidth-optimized kernel density models. *PVLDB* 10, 13 (2017), 2085–2096.
- [33] Andreas Kipf, Thomas Kipf, Bernhard Radke, Viktor Leis, Peter Boncz, and Alfons Kemper. 2019. Learned cardinalities: Estimating correlated joins with deep learning. In *CIDR*.
- [34] Daphne Koller and Nir Friedman. 2009. *Probabilistic graphical models: principles and techniques*. MIT press.
- [35] Frank R Kschischang, Brendan J Frey, and H-A Loeliger. 2001. Factor graphs and the sum-product algorithm. *IEEE Transactions on information theory* 47, 2 (2001), 498–519.
- [36] Seetha Lakshmi and Shaoyu Zhou. 1998. Selectivity estimation in extensible databases-A neural network approach. In *VLDB*, Vol. 98. 24–27.
- [37] Steffen L Lauritzen. 1996. *Graphical models*. Vol. 17. Clarendon Press.
- [38] Viktor Leis, Andrey Gubichev, Atanas Mirchev, Peter Boncz, Alfons Kemper, and Thomas Neumann. 2015. How good are query optimizers, really? *PVLDB* 9, 3 (2015), 204–215.
- [39] Viktor Leis, Bernhard Radke, Andrey Gubichev, Alfons Kemper, and Thomas Neumann. 2017. Cardinality Estimation Done Right: Index-Based Join Sampling. In *CIDR*.
- [40] Feifei Li, Bin Wu, Ke Yi, and Zhuoyue Zhao. 2016. Wander join: Online aggregation via random walks. In *SIGMOD*. 615–629.
- [41] Richard J Lipton, Jeffrey F Naughton, and Donovan A Schneider. 1990. Practical selectivity estimation through adaptive sampling. In *Proceedings of the 1990 ACM SIGMOD international conference on Management of data*. 1–11.
- [42] Jie Liu, Wenqian Dong, Qingqing Zhou, and Dong Li. 2021. Fauce: fast and accurate deep ensembles with uncertainty for cardinality estimation. *Proceedings of the VLDB Endowment* 14, 11 (2021), 1950–1963.
- [43] Qiyu Liu, Yanyan Shen, and Lei Chen. 2021. LHist: Towards Learning Multi-dimensional Histogram for Massive Spatial Data. In *2021 IEEE 37th International Conference on Data Engineering (ICDE)*. IEEE, 1188–1199.
- [44] H-A Loeliger. 2004. An introduction to factor graphs. *IEEE Signal Processing Magazine* 21, 1 (2004), 28–41.
- [45] Lin Ma, Bailu Ding, Sudipto Das, and Adith Swaminathan. 2020. Active learning for ML enhanced database systems. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*. 175–191.
- [46] David JC MacKay, David JC Mac Kay, et al. 2003. *Information theory, inference and learning algorithms*. Cambridge university press.
- [47] Tanu Malik, Randal C Burns, and Nitesh V Chawla. 2007. A Black-Box Approach to Query Cardinality Estimation.. In *CIDR*. Citeseer, 56–67.
- [48] Ryan Marcus, Parimarjan Negi, Hongzi Mao, Nesime Tatbul, Mohammad Alizadeh, and Tim Kraska. 2021. Bao: Learning to steer query optimizers. *SIGMOD* (2021).
- [49] Ryan Marcus, Parimarjan Negi, Hongzi Mao, Chi Zhang, Mohammad Alizadeh, Tim Kraska, Olga Papaemmanouil, and Nesime Tatbul. 2019. Neo: A learned query optimizer. *arXiv preprint arXiv:1904.03711* (2019).
- [50] Ryan Marcus and Olga Papaemmanouil. 2018. Deep reinforcement learning for join order enumeration. In *Proceedings of the First International Workshop on Exploiting Artificial Intelligence Techniques for Data Management*. 1–4.
- [51] Ryan Marcus and Olga Papaemmanouil. 2019. Plan-structured deep neural network models for query performance prediction. *arXiv preprint arXiv:1902.00132* (2019).
- [52] M Muralikrishna and David J DeWitt. 1988. Equi-depth multidimensional histograms. In *Proceedings of the 1988 ACM SIGMOD international conference on Management of data*. 28–36.
- [53] Parimarjan Negi, Matteo Interlandi, Ryan Marcus, Mohammad Alizadeh, Tim Kraska, Marc Friedman, and Alekh Jindal. 2021. Steering query optimizers: A practical take on big data workloads. In *Proceedings of the 2021 International Conference on Management of Data*. 2557–2569.
- [54] Parimarjan Negi, Ryan Marcus, Andreas Kipf, Hongzi Mao, Nesime Tatbul, Tim Kraska, and Mohammad Alizadeh. 2021. Flow-Loss: Learning Cardinality Estimates That Matter. *arXiv preprint arXiv:2101.04964* (2021).
- [55] Hung Q Ngo, Ely Porat, Christopher Ré, and Atri Rudra. 2018. Worst-case optimal join algorithms. *Journal of the ACM (JACM)* 65, 3 (2018), 1–40.
- [56] Yeonsu Park, Seongyun Ko, Sourav S Bhowmick, Kyoungmin Kim, Kijae Hong, and Wook-Shin Han. 2020. G-CARE: A Framework for Performance Benchmarking of Cardinality Estimation Techniques for Subgraph Matching. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*. 1099–1114.
- [57] Hoifung Poon and Pedro Domingos. 2011. Sum-product networks: A new deep architecture. In *ICCV Workshops*. 689–690.
- [58] Viswanath Poosala and Yannis E Ioannidis. 1997. Selectivity estimation without the attribute value independence assumption. In *VLDB*, Vol. 97. 486–495.
- [59] MySQL 8.0 Reference Manual. 2020. Chapter 15.8.10.2 Configuring Non-Persistent Optimizer Statistics Parameters. <https://dev.mysql.com/doc/refman/8.0/en/innodb-statistics-estimation.html> (2020).

- [60] P Griffiths Selinger, Morton M Astrahan, Donald D Chamberlin, Raymond A Lorie, and Thomas G Price. 1979. Access path selection in a relational database management system. In *SIGMOD*. 23–34.
- [61] Utkarsh Srivastava, Peter J Haas, Volker Markl, Marcel Kutsch, and Tam Minh Tran. 2006. Isomer: Consistent histogram construction using query feedback. In *ICDE*. 39–39.
- [62] Michael Stillger, Guy M Lohman, Volker Markl, and Mokhtar Kandil. 2001. LEO-DB2’s learning optimizer. In *PVLDB*, Vol. 1. 19–28.
- [63] Ji Sun and Guoliang Li. 2019. An end-to-end learning-based cost estimator. *VLDB* (2019).
- [64] Kostas Tzoumas, Amol Deshpande, and Christian S Jensen. 2011. Lightweight graphical models for selectivity estimation without independence assumptions. *PVLDB* 4, 11 (2011), 852–863.
- [65] Kostas Tzoumas, Amol Deshpande, and Christian S Jensen. 2013. Efficiently adapting graphical models for selectivity estimation. *Proceedings of the VLDB Endowment* 1, 22 (2013).
- [66] Hai Wang and Kenneth C Sevcik. 2003. A multi-dimensional histogram for selectivity estimation and fast approximate query answering. In *Proceedings of the 2003 conference of the Centre for Advanced Studies on Collaborative research*. 328–342.
- [67] Jiayi Wang, Chengliang Chai, Jiabin Liu, and Guoliang Li. 2021. FACE: a normalizing flow based cardinality estimator. *Proceedings of the VLDB Endowment* 15, 1 (2021), 72–84.
- [68] Chenggang Wu, Alekh Jindal, Saeed Amizadeh, Hiren Patel, Wangchao Le, Shi Qiao, and Sriram Rao. 2018. Towards a learning optimizer for shared clouds. *PVLDB* 12, 3 (2018), 210–222.
- [69] Peizhi Wu and Gao Cong. 2021. A Unified Deep Model of Learning from both Data and Queries for Cardinality Estimation. In *Proceedings of the 2021 ACM SIGMOD International Conference on Management of Data*.
- [70] Ziniu Wu, Parimarjan Negi, Mohammad Alizadeh, Tim Kraska, and Samuel Madden. 2022. FactorJoin: A New Cardinality Estimation Framework for Join Queries. *arXiv preprint arXiv:2212.05526* (2022).
- [71] Ziniu Wu, Amir Shaikhha, Rong Zhu, Kai Zeng, Yuxing Han, and Jingren Zhou. 2020. BayesCard: Revitalizing Bayesian Frameworks for Cardinality Estimation. *arXiv preprint arXiv:2012.14743* (2020).
- [72] Ziniu Wu, Peilun Yang, Pei Yu, Rong Zhu, Yuxing Han, Yaliang Li, Defu Lian, Kai Zeng, and Jingren Zhou. 2022. A Unified Transferable Model for ML-Enhanced DBMS. *CIDR* (2022).
- [73] Ziniu Wu, Rong Zhu, Andreas Pfadler, Yuxing Han, Jiangneng Li, Zhengping Qian, Kai Zeng, and Jingren Zhou. 2020. FSPN: A New Class of Probabilistic Graphical Model. *arXiv preprint arXiv:2011.09020* (2020).
- [74] Zongheng Yang, Amog Kamsetty, Sifei Luan, Eric Liang, Yan Duan, Xi Chen, and Ion Stoica. 2021. NeuroCard: One Cardinality Estimator for All Tables. *PVLDB* 14, 1 (2021), 61–73.
- [75] Zongheng Yang, Eric Liang, Amog Kamsetty, Chenggang Wu, Yan Duan, Xi Chen, Pieter Abbeel, Joseph M Hellerstein, Sanjay Krishnan, and Ion Stoica. 2019. Deep unsupervised cardinality estimation. *PVLDB* (2019).
- [76] Xiang Yu, Guoliang Li, Chengliang Chai, and Nan Tang. 2020. Reinforcement learning with tree-lstm for join order selection. In *2020 IEEE 36th International Conference on Data Engineering (ICDE)*. IEEE, 1297–1308.
- [77] Zhuoyue Zhao, Robert Christensen, Feifei Li, Xiao Hu, and Ke Yi. 2018. Random sampling over joins revisited. In *SIGMOD*. 1525–1539.
- [78] Rong Zhu, Ziniu Wu, Yuxing Han, Kai Zeng, Andreas Pfadler, Zhengping Qian, Jingren Zhou, and Bin Cui. 2021. FLAT: Fast, Lightweight and Accurate Method for Cardinality Estimation. *VLDB* 14, 9 (2021), 1489–1502.

Received April 2022; revised July 2022; accepted August 2022